Development of a High-performance Video Super-resolution Platform using FPGA-CPU Hybrid SoC

A Thesis Topic Proposal
Presented to the Faculty of the
Department of Electronics and Communications Engineering
College of Engineering, De La Salle University

In Partial Fulfillment of
The Requirements for the Degree of
Master of Science in Electronics and Communications Engineering

---

by

Canlas, Reich Rechner D.

April 2015

Abstract

Video super-resolution is the process of reconstructing a high-resolution frame sequence from a low-resolution frame sequence. This problem is ill-posed in nature, that several algorithms have been proposed through the years to increase the quality of the high-resolution output. There are applications which have real-time constraints, such as video stream upscaling and surveillance, therefore, they require high-performance hardware for the super-resolution process. This study aims to design and implement a real-time video super-resolution system that uses the Xilinx Zynq-7000 for a relatively low-cost and low-power yet high-performance super-resolution platform.

Table of Contents

List of Figures

List of Tables

Abbreviations

GPU	**G**raphics **P**rocessing **U**nit
CPU	**C**entral **P**rocessing **U**nit
SR	**S**uper-**R**esolution
PSNR	**P**eak **S**ignal-to-**N**oise **Ratio**
SSIM	**S**tructural **S**imilarity **I**ndex **M**easure
FSIM	**F**eature **S**imilarity **I**ndex **M**easure
FPGA	**F**ield **P**rogrammable **G**ate **A**rray
GPGPU	**G**eneral-**P**urpose computing on **G**raphics **P**rocessing **U**nits
SoC	**S**ystem-**o**n-a-**C**hip
IQA	**I**mage **Q**uality **A**ssessment
HVS	**H**uman **V**ision **S**ystem
HDL	**H**ardware **D**escription **L**anguage

1. INTRODUCTION

1.1 Background of the Study

The search for ever higher screen resolutions led to the advent of ultra high definition television screens and monitors. In the U.S. alone, 80% of households have at least one high-definition screen (Leichtman Research Group, 2015). However, to date, video sources such as terrestrial transmissions here in the Philippines (cite here) do not use the full resolution of modern televisions.

To resolve this problem, a class of methods called "super-resolution" are employed. Super-resolution is the process of rendering or recovering a larger image or video given some low-resolution source (Dong, Loy, & He, 2014). Super-resolution finds its applications in diverse fields of study. Examples include video surveillance, in Caner, a.M. Tekalp, and Heinzelman (2003) and Zhang, Zhang, Shen, and Li (2010), medical imaging in Malczewski and Stasinski (2008), and satellite imaging. Multi-frame image super-resolution methods use a set of LR images to construct an HR image by exploring the spatial correlations in that set (Cheng, Hwang, Jeng, & Lin, 2013). This kind of SR is applicable in laboratory settings. In other cases, single frame SR is more appropriate. These methods try to extract information from only one HR image, making the task much more difficult than multi-frame.

1.2 Statement of the Problem

In many applications such as super-HD (4K) TV, super resolution has to be performed in real time (Shen, Wu, & Deng, 2014). However, as noted by Ishizaka et al. (2013) "it is several times slower than real-time" to upscale videos using commodity hardware. Besides, power consumption is also an important issue. To integrate super-resolution processing into existing systems, there must not be a drastic increase in power footprint. A number of state-of-the-art methods of SR use GPUs and manycore CPUs, which offer a degree of performance at the expense of heavy power consumption and heat dissipation. These current solutions also have unstable frame rates (Wu, Xiang, & Lu, 2011).

We are then confronted by the problem of finding a video super-resolution system that uses a fast algorithm to generate high-quality hi-resolution videos from a low-resolution source while maintaining a relatively small power footprint.

A new class of integrated circuits known as SoCs (Systems-on-a-chip), combining an FPGA and a CPU, are demonstrated to have high performance-per-watt ratios. Thus, they are suitable for the constraints specified in this study.

1.3 Objectives of the Study

1.3.1 General Objective

This study aims to come up with a new video super-resolution algorithm and implement this for use on a hybrid FPGA-CPU (Field Programmable Gate Array - Central Processing Unit) chip.

1.3.2 Specific Objectives

Specifically, the study aims to tackle the following goals:

1. To develop a novel video super-resolution algorithm in terms of PSNR (peak signal-to-noise ratio) and running time.
2. To investigate and determine the performance bottlenecks in the developed video super-resolution algorithm.
3. To investigate the parallelism potential in the algorithm to enhance suitability on the FPGA-CPU hybrid.

1.4 Significance of the Study

The study would contribute to the growing number of literature regarding video SR. It is known that the study of video super-resolution remains in its infancy. High quality and high speed remain difficult to obtain despite the numerous advances in super-resolution through the years. The results

would help in the improvement of picture quality of low-resolution videos on high-resolution monitors (HD, QHD, UHD).

1.5 Scope and Limitations of the Study

Since previous studies (cited in Review of Related Literature) have shown that clear upscaling is practical only at a factor of 4, it has been decided that this study should concern itself with 4x upscaling only.

Two software packages will be used: MATLAB 2015a and Xilinx Vivado Suite 2014.2. The former is an algorithm prototyping environment, while the latter is an HDL (Hardware Description Lanuguage) development environment.

2. REVIEW OF RELATED LITERATURE

2.1 Image and Video Resolutions

Every image and video frame is composed of picture elements, or ”pixels” The HD (High Definition) resolution is 1920x1080 pixels in size

2.2 Image and Video Processing

Image processing is

2.3 Image Quality Assessment (IQA)

Quantification of quality is a prerequisite in the development of any image or video processing algorithm. Currently, there are two primary measures of output image quality, namely, the Peak Signal-to-Noise Ratio (PSNR) and the Structural Similarity Index Measure (SSIM). The choice of PSNR or SSIM is typically arbitrary, with a few informal arguments favoring one or the other (Farsiu, Robinson, Elad, & Milanfar, 2004). To aid in selection of a suitable metric, an analysis of both image metrics is found in Horé and Ziou (2010). They state that a mathematical relationship exists between the two metrics, thus making it possible to predict the PSNR from the SSIM and vice-versa. They only differ in their sensitivity to image degradations as introduced by noise, compression, and hardware limitations.

A recent addition to the list of metrics is the FSIM (Feature Similarity Index Measure) (Zhang, Zhang, Mou, & Zhang, 2011). The metric was proposed on the basis of human visual systems (HVS) understanding an image mainly due to its low-level features, such as edges and zero-crossings.

2.4 Image Super-resolution

Still-image super-resolution (SR) is the reconstruction of a high-resolution (HR) image given one, or a set of, low-resolution (HR) images. In the literature, the words ”super-resolution” and ”upscaling” are typically interchanged, but Takeda, Milanfar, Protter, and Elad (2009) clarified the distinction

between the terms ”upscaling” and ”super-resolution”. They stated that ”if an algorithm which does not receive input frames that are aliased. it will still produce an output with a higher number of pixels and/or frames (i.e., ’upscaled’), but which is not necessarily ’super-resolved’”. Super-resolution began as the problem of image restoration from a noisy signal (Helstrom, 1967). The first known work that directly tackled SR is that of Tsai and Huang (1984).

Traditionally, super-resolution of images is performed with several observed LR images, thus being termed ”multi-frame SR”. This is done in order to remove artifacts introduced by the low-resolution camera sensor (Yang & Huang, 2010). Applications such as medical and satellite imaging prefer the use of multi-frame SR since the minutest of features are crucial to analysis in those fields.

There is another approach which involves only a single observation or image. The limited set of data severely limits the quality obtainable, thus the SR problem becomes ill-posed (Yang & Huang, 2010). Image upscaling for information technology and entertainment is one application that relies on single-frame SR, as there is no available redundancy for the images used in those areas.

Super-resolution is necessary in the following fields of interest:

1. Surveillance video: In Camargo, R.schultz, Wang, Fevig, and He (2010), an SR mosaicking algorithm that stitches and super-resolves UAV (Unmanned Aerial Vehicle) video was implemented on a GPU-CPU pair. They were able to reach as high a PSNR as 41.10 dB for synthetic images. Their application is not real-time, however.

2. Medical imaging: Quan et al. (2010) proposed a real-time algorithm that uses localization-based SR, the superior type of SR microscopy for live cell imaging. The algorithm can achieve between 30-500 fps (frames per second) depending on the speed of the camera used.

3. Video upscaling: This is the main premise of the present study, as with the rest of the papers cited in this chapter.

To the present day, super-resolution remains an active area of research, as evidenced by the wealth of literature cited in this study. The following sections present various approaches to SR that rely on several different models.

2.4.1 Image Observation Model

Several factors affect the output of a digital system, including finite aperture size and finite sensor size (Yang & Huang, 2010). The image observation model, as adopted by common imaging systems, is shown in Figure 2.1.

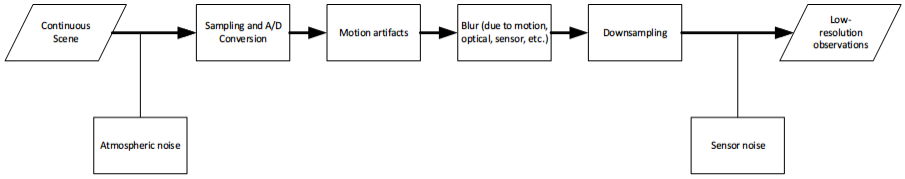

FIGURE 2.1: The Image Observation Model.

2.4.2 Frequency Domain and the Nyquist Theorem

The first SR paper as authored by Tsai and Huang (1984) describes the SR process in the frequency domain. Their algorithm takes advantage of the shift and aliasing properties of the continuous and discrete Fourier transforms, given a set of multiple shifted low resolution images. The main problem is that the typical output of SR methods that depend on Fourier transforms is non-satisfactory for Human Vision Systems. Ultimately, frequency domain methods have been largely superseded by algorithms which take spatial features into account (Yang & Huang, 2010).

Those frequency domain methods primarily rely on the Nyquist theorem (Nyquist, 1928).

2.4.3 Sparse Representation Methods

Digital signal data take up too much storage space, but most of this space does not account for the most significant components of the signal it represents. Compression and alternative representations are therefore required to reduce storage size while preserving fidelity.

A recently-found field of study in signal processing, called compressive sensing Baraniuk, Davenport, Duarte, and Hegde (2011), is posed as a new framework for processing signals.

Dictionary learning is the process of training a set of mutually orthogonal basis vectors in order to create a dictionary matrix. This matrix can then model any signal as a combination of its columns, better known as ”atoms” (Kreutz-Delgado et al., 2003). The goal of using dictionary learning in SR is to find a consistent sparse representation of both the LR and HR patches by training an LR and an HR dictionary. This is also known as the sparse-coding process.

Zeyde, Elad, and Protter (2012) proposed an algorithm that uses the Sparseland model previously developed by Elad and Aharon (2006).

Wright, Huang, Yang, and Ma (2010) jointly trained a dictionary for low resolution and another for high resolution patches to enforce sparse representation similarity for both patch spaces. Their approach is also robust to noise, as it uses local sparse modeling. Yang, Wang, Lin, Cohen, and Huang (2012) similarly stressed the importance of learning two coupled dictionaries (observation dictionary and latent dictionary). However, the difference in their methods is that they used one coupled dictionary learning method for single-image SR.

2.4.4 Computational Intelligence Methods

So far, the previous methods mentioned all have solid mathematical foundations. However, it has been found out that a great number of real-world problems cannot be modeled into well-posed mathematical problems, including super-resolution. A class of algorithms under ”computational intelligence” rely on mimicking natural systems to model and solve such kinds of problems.

Dong et al. (2014) used a deep convolutional neural network in order to learn a mapping between the LR and HR spaces. There are three stages involved: patch extraction and representation, non-linear mapping, and reconstruction. Their method is similar to sparse-coding

2.5 Challenges in Image SR

Researchers still struggle with the following challenges, despite significant advancement in the SR literature.

2.5.1 Image Registration

Image registration is the process of mapping two images both spatially and with respect to intensity (Brown, 1992). Image registration is another ill-posed problem. Artifacts caused by registration problems are more noticeable and annoying than the blurring effect as a result of image interpolation (Yang & Huang, 2010).

2.5.2 Computational Efficiency

According to Yang and Huang (2010), the computational efficiency of SR is severely limited by the fact that there are a large number of unknowns and computationally-expensive matrix manipulations. To alleviate the problem, the author advocates "massive parallel computing". All of the so-called "real-time SR algorithms" can only handle simple motion models, therefore, they cannot be used for real-world videos.

2.5.3 Robustness

SR algorithms are typically sensitive to signal outliers resulting from motion, blur, noise, etc. As the image degradation model parameters are difficult to estimate, researchers nowadays take into account the robustness of their approach (Yang & Huang, 2010).

2.5.4 Edge preservation

It is typical in SR algorithms to lose details or edges in the output image/video. SR techniques for edge preservation have therefore been proposed. Vishnukumar, Nair, and Wilscy (2014) uses self-examples and high-frequency features to provide edge preservation in SR. Their PSNR goes as high as 30.77 dB, SSIM as high as 0.935, and their highest FSIM is 0.955.

2.6 Video super-resolution

Video super-resolution is the extension of image SR to moving pictures. An additional temporal dimension can now be factored in the SR process. It can generally be divided into two categories: incremental and simultaneous (Su, Wu, & Zhou, 2011). The former category is faster but less visually consistent to the human eye. Liu and Sun (2014) mentions that video SR is relatively more challenging than image SR which has been studied for decades, due to the presence of an additional temporal dimension.

2.6.1 Bayesian Methods

Liu and Sun (2014) propose a Bayesian video SR system that can simultaneously estimate the HR frame, motion flow fields, blur kernel, and noise level. Their method works best when the motion is slow and smooth, and would fail if there are significant lighting changes and occlusion. They acknowledged that aliasing both benefits and "derails" super-resolution.

2.6.2 Computational Intelligence Methods

As in image SR, video SR is a highly nonlinear task and is amenable to processing via computational intelligence. For example, Cheng et al. (2013) constructed an artificial neural network (ANN) for video SR. It is a four-stage algorithm, consisting of classifiers to categorize the image, motion-trace volume collection for pixel tracking, temporal adjustment for fast motions and complicated scenes, and ANN learning.

2.7 High Performance Computing Platforms

Yang and Huang (2010) suggest that high-performance hardware matters in tackling super-resolution problems. Typically, image SR algorithms are first developed for computer CPUs. Modern CPUs (central processing units) of computers combine high-frequency processors with a degree of parallelism to add more processing power to algorithms. Even so, the CPU is not enough to handle tasks such as SR in real-time. There are several steps in the SR process that may be implemented as

parallel tasks. Following are the discussions on GPUs, manycore coprocessors, and FPGAs, three parallel platforms commonly in use today.

2.7.1 Graphics Processing Units

GPUs (Graphics Processing Units) have been favored in recent years for this task, as it offers high amounts of parallelism (due to its multiple cores) and compatibility with existing computer systems and programming paradigms. The two major discrete GPU vendors, NVIDIA and AMD, provide specialized tools to offload massively parallel tasks, a process known as GPGPU (General-Purpose computing on GPU).

GPUs are classified as stream processors, because their architecture makes use of a minimal kernel program that processes all data input (the stream). This enables GPUs to process a large amount of data in parallel.

Wu et al. (2011) claims 6x speedup against the same algorithm implemented on a CPU.

2.7.2 Manycore Coprocessors

This class of parallel processors are based off CPU architectures but have more cores than the traditional CPU and are meant to run at a lower frequency. A host CPU passes the appropriate parallel instructions to the manycore coprocessor and subsequently fetches the results of the computation. Manycore processors offer more programmability than GPUs simply by the fact that they share the same architecture as the host CPU. The only known product in this category is that of Intel MIC (Many Integrated Core) architecture (Intel, 2014).

Ishizaka et al. (2013) demonstrated a power-efficient real-time SR system that uses a virtual pipeline to improve the performance as well as the utilization of both the manycore and the host processors. Their set-up was able to achieve 31.5 fps, satisfying the real-time requirement. The downside with their setup is the limited adoption of the MIC platform and the power requirement of about 300 W (Intel, 2014).

2.7.3 Field Programmable Gate Arrays (FPGAs)

FPGAs (Field Programmable Gate Arrays) are logic devices that can be reconfigured by a designer on the field after being manufactured. Since at the lowest level, logic circuits are inherently parallel and real-time, FPGAs offer optimization potential that cannot be realized when using instruction-based platforms such as CPUs and GPUs. FPGAs typically run at much lower frequencies than CPUs and GPUs, making them more power efficient. Higher-end FPGAs even offer the ability to be partially dynamically reconfigured, so that even while it is running, parts of the FPGA fabric gets their design altered (Dye, 2012). These factors makes FPGAs suitable for the most computationally-intensive real-time applications while conserving energy.

Sirowy and Forin (2008) investigated the reasons why an FPGA offers high speedups over sequential processing devices such as CPU, manycore, and GPU. Among these are: the removal of an instruction fetch step, hiding the control instructions, executing multiple instructions in parallel, and pipelining instructions.

2.7.4 Design Considerations and Strategies

Since the SR system of this study is to be integrated into other computing systems, it is imperative to develop an embedded system, one which consumes less energy. The more power used, the more heat is generated. According to Anderson, Dykes, and Riedel (2003), failure rates double for every 15 degree Celsius rise in temperature. In this light, for an embedded system, the GPU and CPU are not applicable processors. Mittal (2014) considered using an ”unconventional core” such as an FPGA to realize lower power-consumption in an embedded system. Struyf, Beugher, Uytsel, Kanters, and Goedem (2014)

2.8 Comparison of CPU, GPU and FPGA

Asano, Maruyama, and Yamaguchi (2009) compared the performance of the CPU, GPU and FPGA in image processing applications. They noted that CPUs are consistently lagging behind the GPU

and FPGA, while the GPU is best for ”naive computation methods” in which processing takes place on a per pixel basis.

Fowers, Brown, Cooke, and Stitt (2012) compared the performance and the energy expended by FPGAs, GPUs and multicore CPUs. This paper is significant to the present study since their focus is on sliding-window algorithms, which take the data on a per-block basis instead of per-pixel. This makes computation more efficient.

Figure **??** illustrates the difference between a sequential processor (CPU) and the FPGA. In the CPU, both data and instructions are being fetched from memory. These instructions are then interpreted into functions to be able to process the data input. Then the intermediate results are then stored in the same memory. On the other hand, the FPGA does not have any instruction to fetch from memory, since all the functions are already defined as hard-wired circuitry. At the same time, the FPGA can have as many parallel functions as possible, and that data may be pushed into the next function stage in just one clock tick.

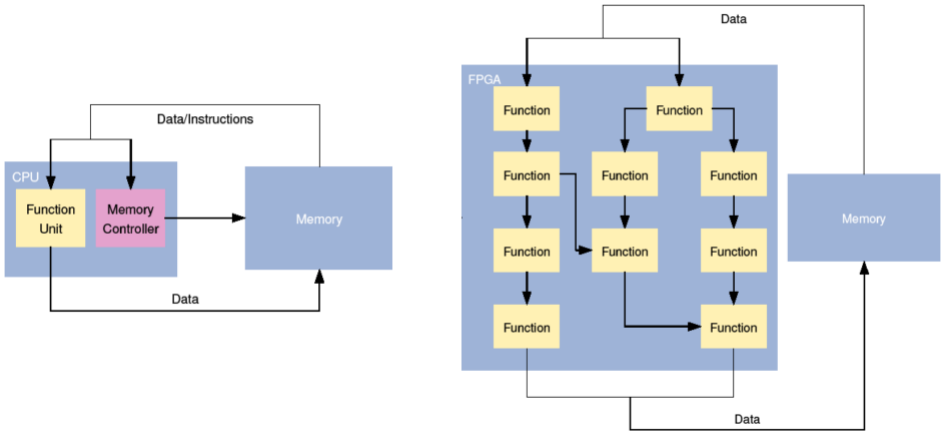

FIGURE 2.2: Comparison of CPU and FPGA working pipelines (Flynn, 2012).

2.8.1 Use in super-resolution applications

The following papers prove the feasibility of an FPGA in SR applications. Angelopoulou, Bouganis, Cheung, and Constantinides (2009) created a real-time video SR system on an FPGA that is robust

against noise. It uses the iterative back projection algorithm. However, the system depends on an adaptive image sensor Szydzik, Callico, and Nunez (2011) constructed a high quality SR system on an FPGA. They were able to achieve 2x upscaling at 25 fps while using only less than 37% of FPGA resources of the state-of-the-art algorithm at that time. Bowen and Bouganis (2008) achieved 3x speedup over equivalent software (CPU) implementations.

2.9 FPGA-CPU System-on-a-chip

The Zynq-7000 programmable System-on-a-Chip is the first-of-its-kind FPGA-CPU hybrid available market. It combines the high-performance, low-power ARM Cortex A9 CPU, high-end computer components such as DDR3 RAM, and a programmable logic fabric similar to that of an FPGA. One advantage of using this SoC is that several peripherals are already built-in and hardwired, eliminating the need for soft IP cores for common devices. Another is that of faster sequential processing. According to Amdahl's Law, the speedup of a parallel program is limited by the time needed for the sequential fraction of the program (Amdahl, 1967). Having both a highly parallel programmable fabric and a fast sequential processor is a great help if an algorithm has to run real-time.

3. METHODOLOGICAL FRAMEWORKS

3.1 Theoretical Considerations

Given a reference image *a* and a test image *b*, both of the same size *M*x*N*,the PSNR is computed using the following equation (Horé & Ziou, 2010):

$$PSNR(a,b) = 20\log_{10}\frac{255}{\sqrt{MSE(a,b)}} \tag{3.1}$$

where

$$MSE(a,b) = \frac{1}{MN}\sum_{i=1}^{M}\sum_{j=1}^{N}(a_{ij}-b_{ij})^2 \tag{3.2}$$

The SSIM is computed using the following equation (Horé & Ziou, 2010):

$$SSIM(a,b) = l(a,b)c(a,b)s(a,b) \tag{3.3}$$

where

$$l(a,b) = \frac{2\mu_a\mu_b + C_1}{\mu_a^2 + \mu_b^2 + C_1} \tag{3.4}$$

$$c(a,b) = \frac{2\sigma_a\sigma_b + C_2}{\sigma_a^2 + \sigma_b^2 + C_2} \tag{3.5}$$

$$s(a,b) = \frac{\sigma_{ab} + C_3}{\sigma_a\sigma_b + C_3} \tag{3.6}$$

Equation 3.4 measures the similarity in luminance and is equal to 1 only if $\mu_a = \mu_b$. Similarly, equation 3.5 compares the standard deviation of the two images (which corresponds to the contrast). It will only equal to 1 if $\sigma_a = \sigma_b$. The structure comparison equation (3.6) measures the correlation between the images using the covariance σ_{ab} between them.

The FSIM is computed using the following equation (Zhang et al., 2011):

$$\frac{\sum_{x\in\Omega} S_L(\mathbf{x}) \cdot PC_m(\mathbf{x})}{\sum_{x\in\Omega} PC_m(\mathbf{x})} \tag{3.7}$$

where Ω is the whole spatial image domain. $S_L(\mathbf{x})$

3.2 Algorithm Framework

Figure 3.1 summarizes the flow of the proposed algorithm. The major components of this algorithm are the deblurring module, the temporal consistency module, and the edge-preservation module. The deblurring module removes image blur caused by motion and the camera sensor. The temporal consistency module ensures that successive video frames are consistently super-resolved with respect to time. It uses the preceding HR frame to accomplish this task. The edge-preservation module takes the finer details of the LR video frames and interpolates the HR version of the edges. A final reconstruction step incorporates output all the three major modules to generate the HR frame sequence.

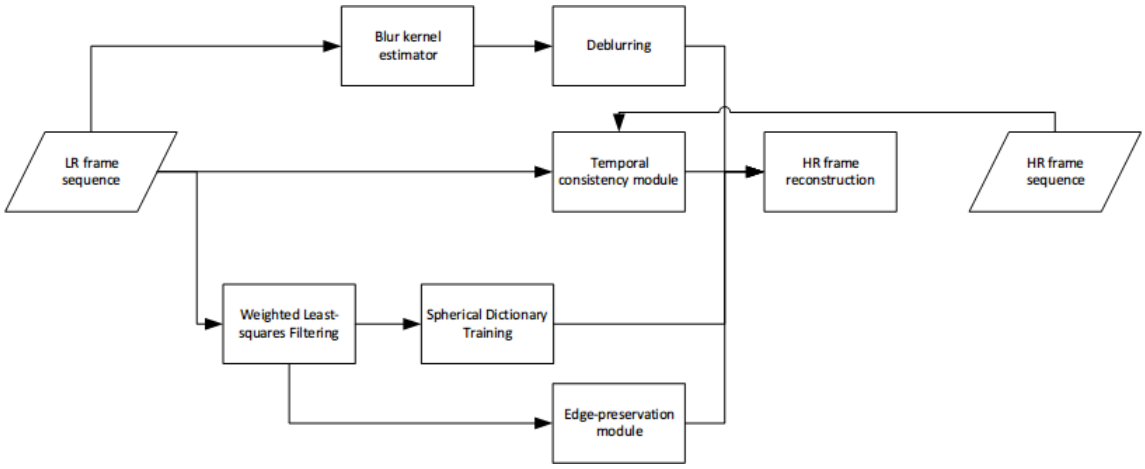

FIGURE 3.1: Framework for the Algorithm.

3.3 Hardware Framework

Figure 3.2 illustrates the flow of data across the hardware devices to be used in the study. A video source such as a camera or prerecorded file will be sent for super-resolution on the SoC, which will then display the result on the monitor in real-time. Initially a conventional computer will serve as the development and evaluation environment for the SR algorithm. A revised version of the algorithm can then be sent to the SoC for further testing and fine-tuning.

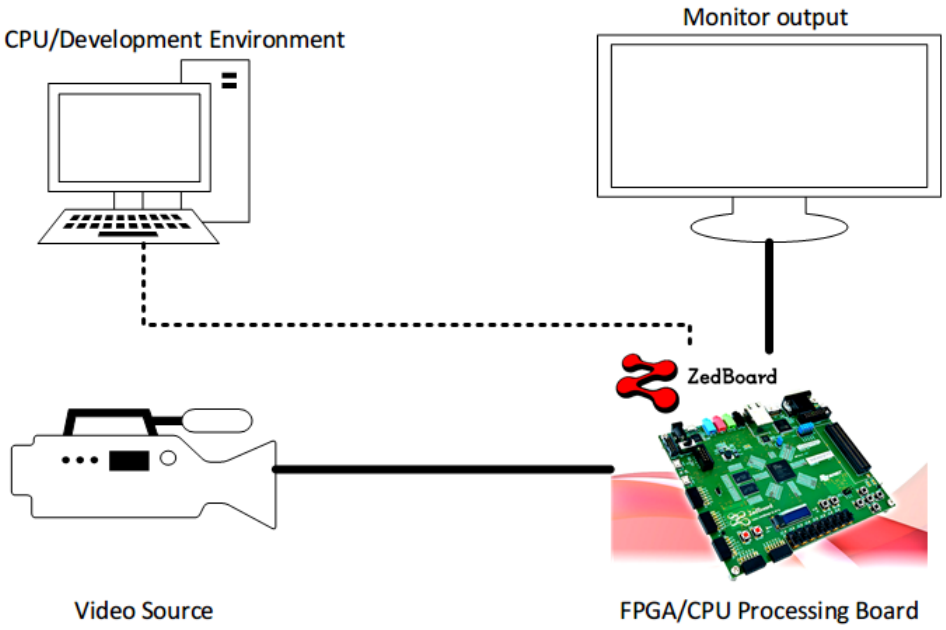

FIGURE 3.2: Framework for Hardware Implementation.

4. METHODOLOGY

4.1 Initial CPU Evaluation

Successive modifications for speed and output quality will be performed until the improvements meet the target +0.1dB increase in quality and 4x speedup. Initially, these measurements will be performed in the CPU using MATLAB functions that measure the PSNR, SSIM, and FSIM.

The initial algorithm runs will take place on a computer with an Intel Core i7 3632QM Processor 2.2 GHz, 16GB RAM, and NVIDIA GeForce GT 730M 2GB. Profiling the algorithm for bottlenecks will be done with the assistance of MATLAB 2015a Profiler. This software package can determine the slowest segments of the algorithm and suggest necessary modifications. MATLAB will also used to code the algorithms and will assist in the preliminary conversion into a form suitable for the Zynq through its Vision HDL Toolbox.

The rationale behind coding a CPU version of the algorithm is to be able to determine which steps in the algorithm have parallelization potential, while the remaining sequential steps could be easily ported over to the CPU segment of the target SoC.

4.2 FPGA Deployment and Comparative Tests

Once the criteria are met in the CPU version of the algorithm, the next step is to parallelize using the available FPGA resources. The platform to be used is the Xilinx Zynq-7000 SoC (embedded CPU+FPGA) on a Digilent Zedboard. The hybrid FPGA and low-power CPU chip enables the dual advantage of quick sequential processing and fine-grained parallel computing. The Xilinx Vivado Design Suite 2014.2 software will facilitate the process of converting the CPU-based algorithm to parallel. In addition, Vivado will be used to fine-tune the algorithm and optimize resource usage in the hardware.

A. Gantt Chart

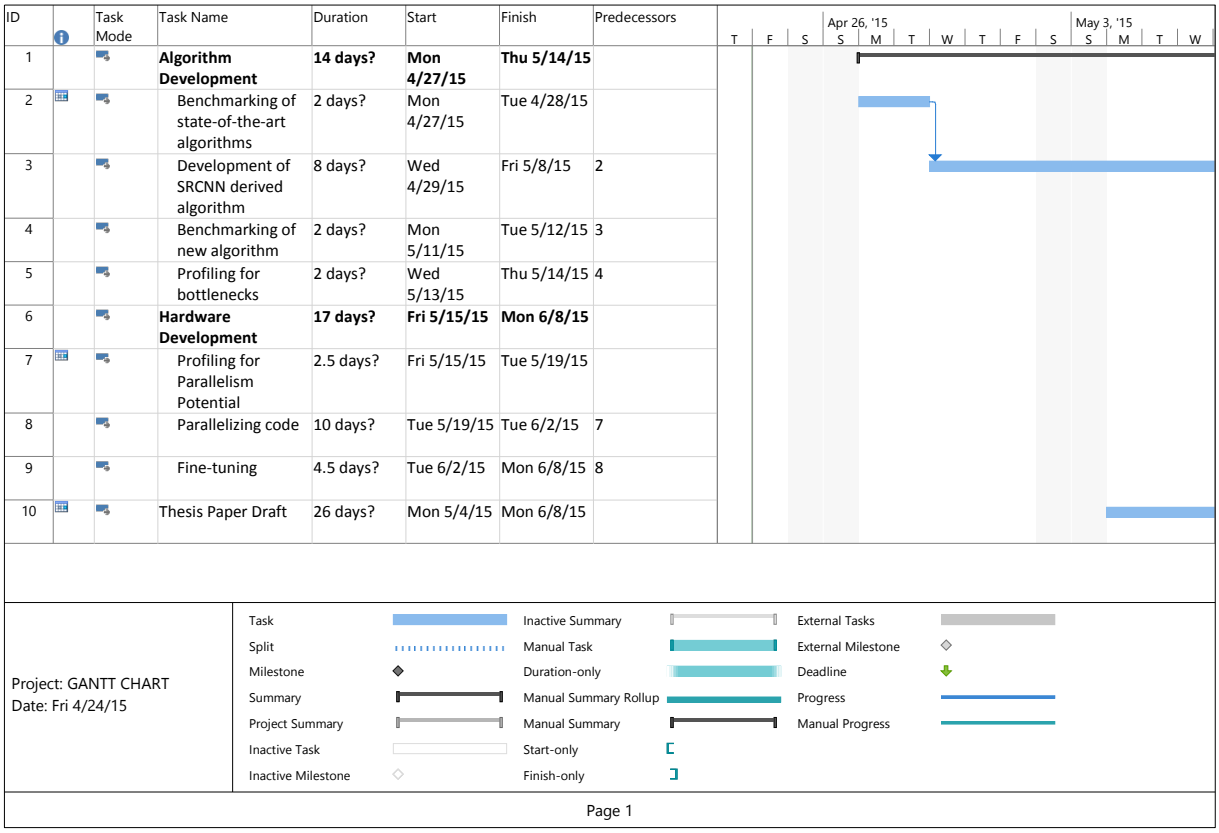

ID	Task Mode	Task Name	Duration	Start	Finish	Predecessors
1		**Algorithm Development**	**14 days?**	**Mon 4/27/15**	**Thu 5/14/15**	
2		Benchmarking of state-of-the-art algorithms	2 days?	Mon 4/27/15	Tue 4/28/15	
3		Development of SRCNN derived algorithm	8 days?	Wed 4/29/15	Fri 5/8/15	2
4		Benchmarking of new algorithm	2 days?	Mon 5/11/15	Tue 5/12/15	3
5		Profiling for bottlenecks	2 days?	Wed 5/13/15	Thu 5/14/15	4
6		**Hardware Development**	**17 days?**	**Fri 5/15/15**	**Mon 6/8/15**	
7		Profiling for Parallelism Potential	2.5 days?	Fri 5/15/15	Tue 5/19/15	
8		Parallelizing code	10 days?	Tue 5/19/15	Tue 6/2/15	7
9		Fine-tuning	4.5 days?	Tue 6/2/15	Mon 6/8/15	8
10		Thesis Paper Draft	26 days?	Mon 5/4/15	Mon 6/8/15	

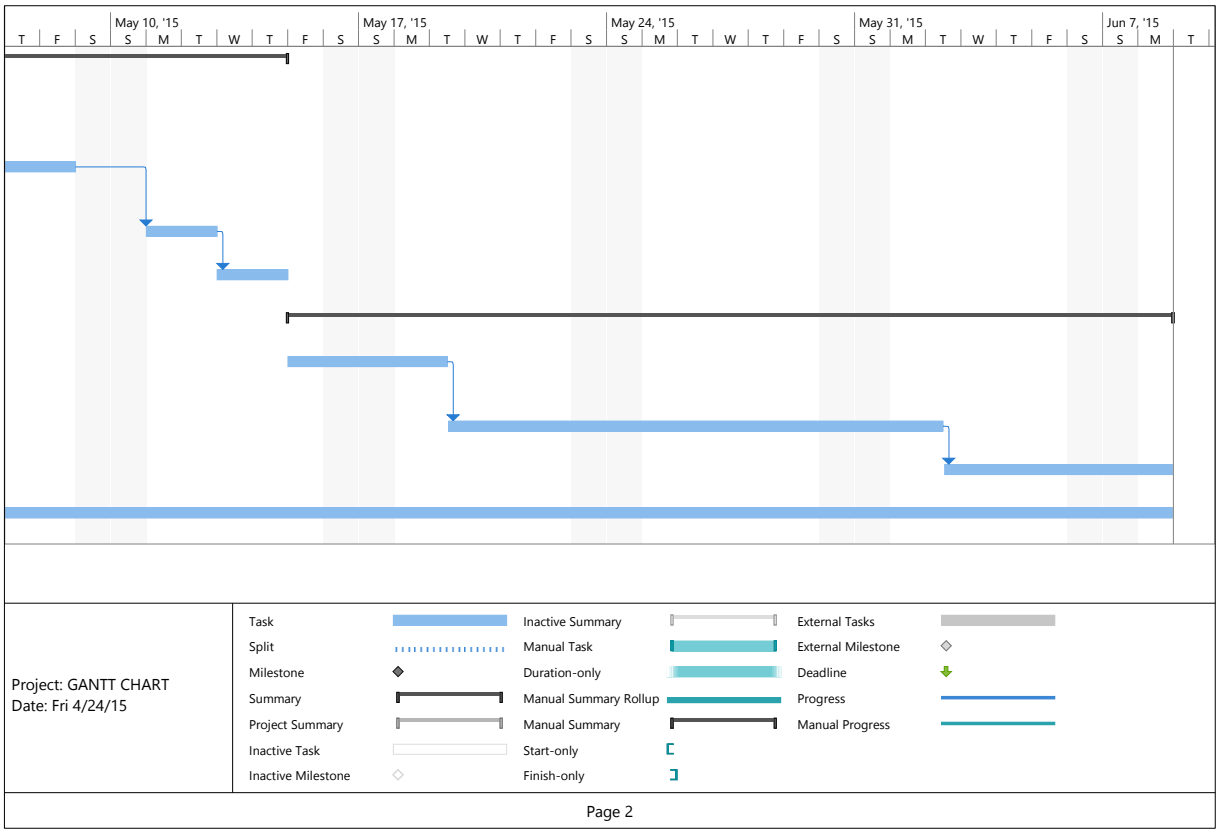
May 10, '15
May 17, '15
May 24, '15
May 31, '15
Jun 7, '15
T F S S M T W T F S S M T W T F S S M T W T F S S M T W T F S S M T
Project: GANTT CHART
Date: Fri 4/24/15
Task
Split
Milestone
Summary
Project Summary
Inactive Task
Inactive Milestone
Inactive Summary
Manual Task
Duration-only
Manual Summary Rollup
Manual Summary
Start-only
Finish-only
External Tasks
External Milestone
Deadline
Progress
Manual Progress
Page 2

B. Proposed Budget

Qty	Particulars	Unit Cost
1	Digilent ZedBoard with Xilinx Zynq-7000 SoC	15000
1	Paper Fees	1000
	Total	16000

References

Amdahl, G. M. (1967). Validity of the single processor approach to achieving large scale computing capabilities. In *Afips spring joint computer conference* (pp. 187–196).

Anderson, D., Dykes, J., & Riedel, E. (2003). More than an interface - SCSI vs. ATA. In *Proceedings of fast '03: 2nd usenix conference on file and storage technologies* (pp. 245–257).

Angelopoulou, M. E., Bouganis, C.-S., Cheung, P. Y. K., & Constantinides, G. a. (2009). Robust Real-Time Super-Resolution on FPGA and an Application to Video Enhancement. *ACM Transactions on Reconfigurable Technology and Systems*, *2*(4), 1–29. doi: 10.1145/1575779.1575782

Asano, S., Maruyama, T., & Yamaguchi, Y. (2009). Performance comparison of FPGA, GPU and CPU in image processing. *FPL 09: 19th International Conference on Field Programmable Logic and Applications*, 126–131. doi: 10.1109/FPL.2009.5272532

Baraniuk, R., Davenport, M. a., Duarte, M. F., & Hegde, C. (2011). An Introduction to Compressive Sensing. *Connexions, Rice University, Houston,Texas*, 118. Retrieved from `http://cnx.org/content/col11133/1.5/`

Bowen, O., & Bouganis, C. (2008). Real-time image super resolution using an FPGA. *Field Programmable Logic and . . .*, 89–94. Retrieved from `http://ieeexplore.ieee.org/xpls/abs_all.jsp?arnumber=4629913`

Brown, L. G. (1992). A survey of image registration techniques. *ACM Computing Surveys*, *24*(4), 325–376. doi: 10.1145/146370.146374

Camargo, A., R.schultz, R., Wang, Y., Fevig, R. a., & He, Q. (2010). GPU-CPU implementation for super-resolution mosaicking of Unmanned Aircraft System (UAS) surveillance video. *Proceedings of the IEEE Southwest Symposium on Image Analysis and Interpretation*, 25–28. doi: 10.1109/SSIAI.2010.5483926

Caner, G., a.M. Tekalp, & Heinzelman, W. (2003). Super resolution recovery for multi-camera surveillance imaging. *2003 International Conference on Multimedia and Expo. ICME '03.*

Proceedings (Cat. No.03TH8698), *1*. doi: 10.1109/ICME.2003.1220866

Cheng, M. H., Hwang, K. S., Jeng, J. H., & Lin, N. W. (2013). Classification-based video super-resolution using artificial neural networks. *Signal Processing*, *93*(9), 2612–2625. Retrieved from http://dx.doi.org/10.1016/j.sigpro.2013.02.013 doi: 10.1016/j.sigpro.2013.02.013

Dong, C., Loy, C. C., & He, K. (2014). Image Super-Resolution Using Deep Convolutional Networks. *arXiv preprint*, 1–14.

Dye, D. (2012). *Partial Reconfiguration of Xilinx FPGAs Using ISE Design Suite* (Vol. 374).

Elad, M., & Aharon, M. (2006). Image denoising via learned dictionaries and sparse representation. *Proceedings of the IEEE Computer Society Conference on Computer Vision and Pattern Recognition*, *1*, 895–900. doi: 10.1109/CVPR.2006.142

Farsiu, S., Robinson, D., Elad, M., & Milanfar, P. (2004). Advances and challenges in super-resolution. *International Journal of Imaging Systems and Technology*, *14*(2), 47–57. doi: 10.1002/ima.20007

Flynn, M. J. (2012). *Using FPGAs for HPC* acceleration: now and in 20 years* (Tech. Rep.). Retrieved from http://tcfpga.org/fpga2012/

Fowers, J., Brown, G., Cooke, P., & Stitt, G. (2012). A performance and energy comparison of FPGAs, GPUs, and multicores for sliding-window applications. *Proceedings of the ACM/SIGDA international symposium on Field Programmable Gate Arrays - FPGA '12*, 47. Retrieved from http://dl.acm.org/citation.cfm?doid=2145694.2145704 doi: 10.1145/2145694.2145704

Helstrom, C. W. (1967). Image Restoration by the Method of Least Squares. *Journal of the Optical Society of America*, *57*(3), 297. doi: 10.1364/JOSA.57.000297

Horé, A., & Ziou, D. (2010). Image quality metrics: PSNR vs. SSIM. *Proceedings - International Conference on Pattern Recognition*, 2366–2369. doi: 10.1109/ICPR.2010.579

Intel. (2014). *Intel Xeon Phi ™ Coprocessor - Datasheet* (Tech. Rep. No. April). Retrieved from http://www.colfax-intl.com/nd/downloads/Xeon-Phi-Coprocessor-Datasheet.pdf

Ishizaka, K., Miyamoto, T., Akimoto, S., Iketani, a., Hosomi, T., & Sakai, J. (2013). Power efficient realtime super resolution by virtual pipeline technique on a server with manycore coprocessors. *IEEE Symposium on Low-Power and High-Speed Chips - Proceedings for 2013 COOL Chips XVI*, *2*, 2–4. doi: 10.1109/CoolChips.2013.6547918

Kreutz-Delgado, K., Murray, J. F., Rao, B. D., Engan, K., Lee, T.-W., & Sejnowski, T. J. (2003). Dictionary learning algorithms for sparse representation. *Neural computation*, *15*(2), 349–396. doi: 10.1162/089976603762552951

Leichtman Research Group. (2015). *HDTV SETS NOW IN OVER 80% OF U. S. HOUSEHOLDS* (Tech. Rep.). Durham.

Liu, C., & Sun, D. (2014). On bayesian adaptive video super resolution. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, *36*(2), 346–360. doi: 10.1109/TPAMI.2013.127

Malczewski, K., & Stasinski, R. (2008). Toeplitz-based iterative image fusion scheme for MRI. *Proceedings - International Conference on Image Processing, ICIP*, 341–344. doi: 10.1109/ICIP.2008.4711761

Mittal, S. (2014). A Survey of Techniques For Improving Energy Efficiency in Embedded Computing Systems. *International Journal of Computer Aided Engineering and Technology*. Retrieved from `http://arxiv.org/abs/1401.0765` doi: 10.1504/IJCAET.2014.065419

Nyquist, H. (1928). Certain topics in telegraph transmission theory. *Proceedings of the IEEE*, *90*(2), 280–305. doi: 10.1109/5.989875

Quan, T., Li, P., Long, F., Zeng, S., Luo, Q., Hedde, P. N., ... Huang, Z.-L. (2010). Ultra-fast, high-precision image analysis for localization-based super resolution microscopy. *Optics express*, *18*(11), 11867–11876. doi: 10.1364/OE.18.011867

Shen, Y., Wu, X., & Deng, X. (2014, December). GPU-aided real-time image/video super resolution based on error feedback. In *2014 ieee visual communications and image processing conference* (pp. 286–290). IEEE. Retrieved from `http://ieeexplore.ieee.org/lpdocs/epic03/wrapper.htm?arnumber=7051560` doi: 10.1109/VCIP.2014.7051560

Sirowy, S., & Forin, A. (2008). Where's the Beef ? Why FPGAs Are So Fast. *Microsoft*

Research(September).

Struyf, L., Beugher, S. D., Uytsel, D. H. V., Kanters, F., & Goedem, T. (2014). The battle of the giants : a case study of GPU vs FPGA optimisation for real-time image processing. In *International conference on pervasive and embedded computing and communication systems – peccs 2014* (pp. 112–119). Lisbon, Portugal: VISIGRAPP. Retrieved from http://www.peccs.org/?y=2014

Su, H., Wu, Y., & Zhou, J. (2011). Adaptive incremental video super-resolution with temporal consistency. *Proceedings - International Conference on Image Processing, ICIP*, 1149–1152. doi: 10.1109/ICIP.2011.6115632

Szydzik, T., Callico, G. M., & Nunez, A. (2011). Efficient FPGA Implementation of a High-Quality Super- Resolution Algorithm with Real-Time Performance. , *57*(2), 664–672.

Takeda, H., Milanfar, P., Protter, M., & Elad, M. (2009). Super-resolution without explicit subpixel motion estimation. *IEEE Transactions on Image Processing*, *18*(9), 1958–1975. doi: 10.1109/TIP.2009.2023703

Tsai, R., & Huang, T. S. (1984). Multiframe image restoration and registration. *Advances in computer vision and Image Processing*, *1*(2), 317–339.

Vishnukumar, S., Nair, M. S., & Wilscy, M. (2014). Edge preserving single image super-resolution with improved visual quality. *Signal Processing*, *105*, 283–297. doi: 10.1016/j.sigpro.2014.05.033

Wright, J., Huang, T., Yang, J., & Ma, Y. (2010). Image Super-Resolution via Sparse Representation. *IEEE transactions on image processing : a publication of the IEEE Signal Processing Society*, 1–13. Retrieved from http://www.ncbi.nlm.nih.gov/pubmed/20483687$\delimiter"026E30F$nhttp://ieeexplore.ieee.org/lpdocs/epic03/wrapper.htm?arnumber=4587647 doi: 10.1109/TIP.2010.2050625

Wu, X., Xiang, H., & Lu, P. (2011). A GPU Accelerated Algorithm for Compressive Sensing Based Image Super-Resolution. *2011 Workshop on Digital Media and Digital Content Management*(60633030), 198–202. doi: 10.1109/DMDCM.2011.10

Yang, J., & Huang, T. (2010). Image super-resolution: Historical overview and future challenges.

In *Super-resolution imaging* (pp. 3–25). Retrieved from http://books.google.com/books?hl=en&lr=&id=fjTUbMnvOkgC&oi=fnd&pg=PA1&dq=Image+super-resolution:+Historical+overview+and+future+challenges&ots=53GZConOGy&sig=oHgGJVo_57Tv19cxLV_XMtC1OPw

Yang, J., Wang, Z., Lin, Z., Cohen, S., & Huang, T. (2012). Coupled dictionary training for image super-resolution. *IEEE Transactions on Image Processing*, *21*(8), 3467–3478. doi: 10.1109/TIP.2012.2192127

Zeyde, R., Elad, M., & Protter, M. (2012). On single image scale-up using sparse-representations. *Lecture Notes in Computer Science (including subseries Lecture Notes in Artificial Intelligence and Lecture Notes in Bioinformatics)*, *6920 LNCS*(1), 711–730. doi: 10.1007/978-3-642-27413-8_47

Zhang, L., Zhang, H., Shen, H., & Li, P. (2010). A super-resolution reconstruction algorithm for surveillance images. *Signal Processing*, *90*(3), 848–859. Retrieved from http://dx.doi.org/10.1016/j.sigpro.2009.09.002 doi: 10.1016/j.sigpro.2009.09.002

Zhang, L., Zhang, L., Mou, X., & Zhang, D. (2011). FSIM : A Feature Similarity Index for Image Quality Assessment. *IEEE Transactions on Image Processing*, *20*(8), 2378–2386.